

Manike B2B AI Engine

Project Workflow Documentation

Multi-Tenant AI Scene Orchestrator & Itinerary Engine

Version: 2.0.0

Architecture: FastAPI + PostgreSQL + Milvus

Date: February 2026

1. Architecture Overview

The Manike B2B AI Engine is a multi-tenant backend system designed for travel companies. It uses a hybrid database architecture combining PostgreSQL for structured data and Milvus for vector-based semantic search. The system orchestrates multiple AI services to generate cinematic travel content including images, videos, and itineraries.

Technology Stack

- Framework: FastAPI (Python)
- Relational DB: PostgreSQL (via SQLAlchemy/SQLModel)
- Vector DB: Milvus (for AI-powered semantic search)
- Authentication: JWT (JSON Web Tokens)
- Media Processing: FFmpeg
- Storage: S3-compatible object storage
- AI Models: LLM (GPT-4/Claude), Image Gen, Video Gen (pluggable)

Project Structure

```
menike_backend_poc/
|-- app/
|   |-- api/           # API endpoint routers
|   |   |-- auth.py     # Login & JWT token generation
|   |   |-- scenes.py   # Scene orchestration endpoints
|   |   |-- itinerary.py # Itinerary generation endpoints
|   |   |-- experiences.py # Experience CRUD (Milvus)
|   |   |-- tenants.py   # Tenant management (Milvus)
|   |-- core/           # Core infrastructure
|   |   |-- auth.py      # JWT logic & password hashing
|   |   |-- database.py  # PostgreSQL engine & sessions
|   |   |-- milvus_client.py # Milvus connection & queries
|   |-- models/          # Data models
|   |   |-- sql_models.py # Tenant, User, Scene (PostgreSQL)
|   |   |-- milvus_schema.py # Vector collection schemas
|   |-- services/         # Business logic
|   |   |-- orchestrator.py # AI scene pipeline
|   |   |-- generators.py  # LLM, Image, Video generators
|   |   |-- media_processor.py # FFmpeg wrapper
|   |   |-- storage.py     # S3 file upload
|   |-- main.py          # FastAPI app entry point
|-- seed_db.py          # Database seeding script
-- requirements.txt      # Python dependencies
-- .env                  # Environment variables
```

2. Complete Working Flow

The following describes the end-to-end flow of how a request travels through the system, from user authentication to AI content generation and storage.

Phase A: Authentication Flow

1 User Login

The client sends a POST request to /auth/login with email and password. The system verifies the credentials against the PostgreSQL 'user' table using PBKDF2_SHA256 password hashing.

2 JWT Token Generation

On successful authentication, the server generates a JWT token containing the user's ID and tenant_id. This token is valid for 24 hours and must be included in the Authorization header of all subsequent requests.

3 Tenant Extraction

Every protected endpoint uses the get_current_tenant_id dependency, which decodes the JWT token, retrieves the user from PostgreSQL, and returns their tenant_id. This ensures complete data isolation between tenants.

Phase B: Scene Orchestration Flow (Core AI Pipeline)

4 Scene Request

The authenticated client sends a POST request to /scenes/ with a scene name and description (e.g., 'Galle Fort Sunset' with a cinematic description).

5 Database Record Creation

The SceneOrchestrator creates a new Scene record in PostgreSQL with status='pending', linked to the authenticated tenant_id. It then updates the status to 'processing'.

6 LLM Prompt Engineering

The LLMPromptEngine analyzes the scene description and generates structured prompts:

- image_prompt: Optimized for photorealistic image generation
- video_prompt: Optimized for cinematic video generation
- negative_prompt: Quality guard rails for AI models

7 AI Image Generation

The ImageGenerator receives the image_prompt and produces a high-quality travel image. Currently uses a mock implementation; in production, this connects to Stable Diffusion, DALL-E, or Leonardo AI.

8 AI Video Generation

The VideoGenerator receives the video_prompt and produces a cinematic video clip. Currently uses a mock; in production, connects to Veo 3 or Runway ML.

9 Media Processing (FFmpeg)

The MediaProcessor uses FFmpeg to optimize the generated video:

- Compression for web streaming
- Format standardization (MP4/H.264)
- Resolution and bitrate optimization

10 Storage Upload

The optimized media file is uploaded to S3-compatible storage. A public URL is generated and stored back in the Scene record in PostgreSQL.

11 Completion

The Scene record is updated with status='completed', the media_url, and a timestamp. The complete scene object is returned to the client.

Phase C: Itinerary Generation Flow

12 Itinerary Request

The client sends a POST to /itinerary/generate with destination, number of days, and interests. The system generates a structured multi-day travel plan.

13 AI Planning

The LLM creates a logical itinerary with activities, travel times, and distances. It also generates image prompts for each activity for visual enrichment.

Phase D: Experience Vector Search Flow

14 Experience Storage


Travel experiences are stored in Milvus as vector embeddings (768-dimensional). Each experience is tagged with a tenant_id for data isolation.

15 Semantic Search

The client can search for similar experiences using POST /experiences/search/ with a query vector. Milvus performs HNSW-based approximate nearest neighbor search filtered by tenant_id, returning the most relevant results ranked by cosine similarity.

3. Data Flow Diagram

The following illustrates how data moves through the system layers:



4. API Endpoints Summary

Method	Endpoint	Description
POST	/auth/login	Authenticate user, returns JWT token
POST	/scenes/	Create a new AI-orchestrated scene
GET	/scenes/	List all scenes for the authenticated tenant
POST	/itinerary/generate	Generate a multi-day travel itinerary
GET	/experiences/	List all experiences from Milvus
POST	/experiences/	Create a new experience in Milvus
POST	/experiences/search/	Semantic vector search for experiences
POST	/tenants/	Create a new tenant in Milvus
GET	/tenants/	List all tenants
GET	/tenants/{id}	Get a specific tenant by ID
PUT	/tenants/{id}	Update a tenant
DELETE	/tenants/{id}	Delete a tenant

5. Database Models

PostgreSQL Tables

- Tenant: id, name, api_key, config, created_at
- User: id, tenant_id (FK), email, hashed_password, role, created_at
- Scene: id, tenant_id, name, description, status, media_url, created_at, updated_at

Milvus Collections

- experiences: id, tenant_id, embedding (768-dim), metadata (JSON), slug
- tenants: id, name, apikey, metadata (JSON), embedding (2-dim placeholder)

6. How to Run the Project

Prerequisites

- Python 3.10+
- Docker Desktop (for PostgreSQL and Milvus)
- pip (Python package manager)

Step 1: Start Docker Services

```
# Start PostgreSQL
docker run -d --name postgres_local \
-e POSTGRES_USER=admin \
-e POSTGRES_PASSWORD=admin \
-e POSTGRES_DB=menike \
-p 5432:5432 postgres:latest

# Start Milvus (using docker-compose)
# Follow: https://milvus.io/docs/install_standalone-docker.md
```

Step 2: Configure Environment

```
# .env file contents:
DATABASE_URL=postgresql://admin:admin@localhost:5432/menike
```

Step 3: Install Dependencies & Seed

```
pip install -r requirements.txt
python seed_db.py
```

Step 4: Run the Server

```
export PYTHONPATH=$PYTHONPATH:$(pwd)
python app/main.py

# Server starts at http://localhost:8000
# Swagger UI at http://localhost:8000/docs
```

Step 5: Test the Flow via Swagger

1. Open <http://localhost:8000/docs>
2. POST /auth/login with: username=admin@manike.ai, password=admin123
3. Copy the access_token from the response
4. Click 'Authorize' button, paste the token
5. POST /scenes/ to create a scene
6. GET /scenes/ to verify the scene was created